

A large blue square containing a large white number '2'. A thin vertical line extends from the bottom of the square towards the word 'Solutions'.

# 2

## Solutions



**2.1** `addi $s0, $s2, -5`  
`add $s0, $s0, $s1`  
`[addi f,h,-5 (note, no subi) add f,f,g]`

**2.2** `f = g+h+i`

**2.3** `sub $t0, $s3, $s4`  
`sll $t0, $t0, 2`  
`add $t0, $s6, $t0`  
`lw $t1, 0($t0)`  
`sw $t1, 32($s7)`

**2.4** `B[g] = A[f] + A[f+1]`  
`sll $t0, $s0, 2` # `t0 = f * 4`  
`add $t0, $s6, $t0` # `t0 = &A[f]`  
`lw $t1, 0($t0)` # `s0 = A[f]`  
`lw $t2, 4($t0)` # `s1 = A[1+f]`  
`add $t0, $t1, $t2` # `t0 = A[f] + A[1+f]`  
`sll $t1, $s1, 2` # `t1 = g * 4`  
`add $t1, $s7, $t1` # `t1 = &B[g]`  
`sw $t0, 0($t1)` # `B[g] = A[f] + A[1+f]`

## 2.5

Little-Endian		Big-Endian	
Address	Data	Address	Data
3	ab	3	12
2	cd	2	ef
1	ef	1	cd
0	12	0	ab

## 2.6 2882400018

**2.7** `sll $t0, $s3, 2` # `$t0 <-- 4*i`  
`add $t0, $t0, $s6` # `$t0 <-- Addr(A[i])`  
`lw $t0, 0($t0)` # `$t0 <-- A[i]`  
`sll $t1, $s4, 2` # `$t1 <-- 4*j`  
`add $t1, $t1, $s6` # `$t1 <-- Addr(A[j])`  
`lw $t1, 0($t1)` # `$t1 <-- A[j]`  
`add $t0, $t0, $t1` # `t1 <-- A[i] + A[j]`  
`sw $t0, 32($s7)` # `B[8] <-- A[i] + A[j]`

**2.8**  $f = 2 * (\&A)$

```

addi $t0, $s6, 4
add  $t1, $s6, $0
sw   $t1, 0($t0)
lw   $t0, 0($t0)
add  $s0, $t1, $t0

```

**2.9**

	type	opcode	rs	rt	rd	immed
addi \$t0, \$s6, 4	I-type	8	22	8		4
add \$t1, \$s6, \$0	I-type	0	22	0	9	
sw \$t1, 0(\$t0)	I-type	43	8	9		0
lw \$t0, 0(\$t0)	I-type	35	8	8		0
add \$s0, \$t1, \$t0	R-type	0	9	8	16	

	type	opcode, funct3,7	rs1	rs2	rd	Imm
addi x30, X 10.8	I-type	0x13, 0x0, -	10	-	30	8
addi x31, X 10.0	R-type	0x13, 0x0, -	10	-	31	0
sd x31, 0(x30)	S-type	0x23, 0x3, -	31	30	-	0
ld x30, 0(x30)	I-type	0x3, 0x3, -	30	-	30	0
add x5, X30, x31	R-type	0x33, 0x0, 0x0	30	31	5	-

**2.10**

**2.10.1** 0x50000000

**2.10.2** overflow

**2.10.3** 0xB0000000

**2.10.4** no overflow

**2.10.5** 0xD0000000

**2.10.6** overflow

**2.11**

**2.11.1** There is an overflow if  $128 + \$s1 > 2^{31} - 1$ .

In other words, if  $\$s1 > 2^{31} - 129$ .

There is also an overflow if  $128 + \$s1 < -2^{31}$ .

In other words, if  $\$s1 < -2^{31} - 128$  (which is impossible given the range of  $\$s1$ ).

**2.11.2** There is an overflow if  $128 - \$s1 > 2^{31} - 1$ .

In other words, if  $\$s1 < -2^{31} + 129$ .

There is also an overflow if  $128 - \$s1 < -2^{31}$ .

In other words, if  $\$s1 > 2^{31} + 128$  (which is impossible given the range of  $\$s1$ ).

**2.11.3** There is an overflow if  $\$s1 - 128 > 2^{31} - 1$ .  
 In other words, if  $\$s1 > 2^{31} + 127$  (which is impossible given the range of  $\$s1$ ).  
 There is also an overflow if  $\$s1 - 128 < -2^{31}$ .  
 In other words, if  $\$s1 < -2^{31} + 128$ .

**2.12** R-type: `add $s0, $s0, $s0`

**2.13** I-type: `0xAD490020`

**2.14** R-type: `sub $v1, $v1, $v0, 0x00621822`

**2.15** I-type: `lw $v0, 4($at), 0x8C220004`

## 2.16

**2.16.1** The opcode would expand from 7 bits to 9.  
 The `rs1`, `rs2`, and `rd` fields would increase from 5 bits to 7 bits.

**2.16.2** The opcode would expand from 7 bits to 9.  
 The `rs1` and `rd` fields would increase from 5 bits to 7 bits. This change does not affect the `imm` field *per se*, but it might force the ISA designer to consider shortening the immediate field to avoid an increase in overall instruction size.

**2.16.3** \* Increasing the size of each bit field potentially makes each instruction longer, potentially increasing the code size overall.  
 \* However, increasing the number of registers could lead to less register spillage, which would reduce the total number of instructions, possibly reducing the code size overall.

## 2.17

**2.17.1** `0xBABEFEF8`

**2.17.2** `0xAAAAAAAA0`

**2.17.3** `0x00005545`

**2.18** It can be done in eight MIPS instructions:

```

srl  $t0, $t0, 11
sll  $t0, $t0, 26
ori  $t2, $0, 0x03ff
sll  $t2, $t2, 16
ori  $t2, $t2, 0xffff
and  $t1, $t1, $t2
or   $t1, $t1, $t0
    
```

**2.19** `nor $t1, $t2, $t2`

**2.20** `lw $t3, 0($s1)`  
`sll $t1, $t3, 4`

**2.21** `$t2 = 3`

**2.22**

**2.22.1** `[0x20000000, 2FFFFFFC]`

**2.22.2** `[1FFE0000, 2001FFFC]`

**2.23**

**2.23.1** The I-type instruction format would be most appropriate because it would allow the branch target to be encoded into the immediate field and the register address to be encoded into the `rt` field, whose datapath supports both read and write access the register file.

**2.23.2** It can be done in three instructions:

Loop:

```
addi    $s0,$s0,-1    # subtract 1 from $s0
bltz    $s0,loop      # continue if $s0 not
                    negative
addi    $s0,$s0,1      # add back 1 that shouldn't
                    have been subtracted
```

**2.24**

**2.24.1** The final value of `$s2` is 20.

**2.24.2** `i = 10;`  
`do {`  
 `B += 2;`  
 `i = i - 1;`  
`} while (i > 0)`

**2.24.3** `5*N` instructions.

**2.25** The C code can be implemented in RISC-V assembly as follows.

```

        addi    $t0,    $0, 0
        beq     $0,     $0, TEST1
LOOP1:  addi    $t1,    $0, 0
        beq     $0,     $0, TEST2
LOOP2:  add     $t3,    $t0, $t1
        sll     $t2,    $t1, 4
        add     $t2,    $t2, $s2
        sw     $t3,    ($t2)
        addi    $t1,    $t1, 1
TEST2:  slt     $t2,    $t1, $s1
        bne     $t2,    $0, LOOP2
        addi    $t0,    $t0, 1
TEST1:  slt     $t2,    $t0, $s0
        bne     $t2,    $0, LOOP1
    
```

**2.26** Answers will vary, but should require approximately 14 MIPS instructions, and when  $a = 10$  and  $b = 1$ , should result in approximately 158 instructions being executed.

**2.27** // Assume MemArray is an array of 32-bit words.  
 for (i=0; i<100; i++) {  
 result += MemArray[i];  
 }  
 return result;

**2.28**

```

        addi    $t1,    $s0, 400
LOOP:   lw      $s1,    0($t1)
        add     $s2,    $s2, $s1
        addi    $t1,    $t1, -4
        bne     $t1,    $s0, LOOP
    
```

```

2.29 fib:   addi $sp,$sp, -12      # make room on stack
            sw   $ra, 8($sp)      # push $ra
            sw   $s0, 4($sp)      # push $s0
            sw   $a0, 0($sp)      # push $a0 (N)
            bgt  $a0,$0, test2     # if n>0, test if n=1
            add  $v0, $0, $0      # else fib(0) = 0
            j    rtn              #
test2:      addi $t0, $0, 1        #
            bne  $t0, $a0, gen     # if n>1, gen
            add  $v0, $0, $t0      # else fib(1) = 1
            j    rtn              #
gen:        subi $a0,$a0,1         # n-1
            jal  fib               # call fib(n-1)
            add  $s0,$v0, $0        # copy fib(n-1)
            sub  $a0,$a0,1         # n-2
            jal  fib               # call fib(n-2)
            add  $v0, $v0, $s0      # fib(n-1)+fib(n-2)
rtn:        lw   $a0, 0($sp)        # pop $a0
            lw   $s0, 4($sp)        # pop $s0
            lw   $ra, 8($sp)        # pop $ra
            addi $sp, $sp, 12       # restore sp
            jr   $ra
# fib(0) = 12 instructions, fib(1) = 14 instructions,
# fib(N) = 26 + 18N instructions for N >=2
done:
jalr x0, x1

```

**2.30** after calling function fib:

```

old $sp -> 0x7ffffffc   ???
           -4           contents of register $ra for
                        fib(N)
           -8           contents of register $s0 for
                        fib(N)
$sp->      -12          contents of register $a0 for
                        fib(N)
there will be N-1 copies of $ra, $s0 and $a0

```

```

2.31 f: addi    $sp,$sp,-12
        sw      $ra,8($sp)
        sw      $s1,4($sp)
        sw      $s0,0($sp)
        move    $s1,$a2
        move    $s0,$a3
        jal     func
        move    $a0,$v0
        add     $a1,$s0,$s1
        jal     func
        lw      $ra,8($sp)
        lw      $s1,4($sp)
        lw      $s0,0($sp)
        addi    $sp,$sp,12
        jr      $ra
    
```

**2.32** We can use the tail-call optimization for the second call to func, but then we must restore \$ra, \$s0, \$s1, and \$sp before that call. We save only one instruction (jr \$ra).

**2.33** Register \$ra is equal to the return address in the caller function, registers \$sp and \$s3 have the same values they had when function f was called, and register \$t5 can have an arbitrary value. For register \$t5, note that although our function f does not modify it, function func is allowed to modify it so we cannot assume anything about the of \$t5 after function func has been called.

```

2.34 MAIN: addi    $sp, $sp, -4
        sw      $ra, ($sp)
        add     $t6, $0, 0x30 # '0'
        add     $t7, $0, 0x39 # '9'
        add     $s0, $0, $0
        add     $t0, $a0, $0
LOOP:   lb      $t1, ($t0)
        slt     $t2, $t1, $t6
        bne     $t2, $0, DONE
        slt     $t2, $t1, $t7
        bne     $t2, $0, DONE
        sub     $t1, $t1, $t6
        beq     $s0, $0, FIRST
        mul     $s0, $s0, 10
FIRST:  add     $s0, $s0, $t1
        addi    $t0, $t0, 1
        j       LOOP
DONE:   add     $v0, $s0, $0
        lw      $ra, ($sp)
        addi    $sp, $sp, 4
        jr      $ra
    
```





## 2.35

**2.35.1** 0x11

**2.35.2** 0x44

**2.36** Generally, all solutions are similar:

```
lui $t1, top_16_bits
ori $t1, $t1, bottom_16_bits
```

**2.37** setmax:

```
ll $t0,0($a0) # $t0 = *shvar
sub $t1,$t0,$a1 # $t1 = *shvar-x
bgez $t1,skip # if result is zero or positive, then
x <= *shvar, so keep original value of *shvar, else
replace with x
add $t0,$0,$a1
```

skip:

```
sc $t0,0($a0) # try to store result to shvar
beqz $t0,setmax # repeat if operation was not atomic
```

**2.38** When two processors A and B begin executing this routine at the same time, at most one of them will execute the store-conditional instruction successfully, while the other will be forced to retry the operation. If processor A's store-conditional succeeds initially, then B will re-enter the try block, and it will see the new value of shvar written by A when it finally succeeds. The hardware guarantees that both processors will eventually execute the code completely.

## 2.39

**2.39.1** No. The resulting machine would be slower overall.

Current CPU requires (num arithmetic \* 1 cycle) + (num load/store \* 10 cycles) + (num branch/jump \* 3 cycles) =  $500 * 10^6 * 1 + 300 * 10^6 * 10 + 100 * 10^6 * 3 = 3800 * 10^6$  cycles.

The new CPU requires  $(.75 * \text{num arithmetic} * 1 \text{ cycle}) + (\text{num load/store} * 10 \text{ cycles}) + (\text{num branch/jump} * 3 \text{ cycles}) = 375 * 10^6 * 1 + 300 * 10^6 * 10 + 100 * 10^6 * 3 = 3675 * 10^6$  cycles.

However, given that each of the new CPU's cycles is 10% longer than the original CPU's cycles, the new CPU's  $3675 * 10^6$  cycles will take as long as  $4042.5 * 10^6$  cycles on the original CPU.

**2.39.2** If we double the performance of arithmetic instructions by reducing their CPI to 0.5, then the CPU will run the reference program in  $(500 * .5) + (300 * 10) + 100 * 3 = 3550$  cycles. This represents a speedup of 1.07.

If we improve the performance of arithmetic instructions by a factor of 10 (reducing their CPI to 0.1), then the CPU will run the reference program in  $(500 * .1) + (300 * 10) + 100 * 3 = 3350$  cycles. This represents a speedup of 1.13.





## 2.40

**2.40.1** Take the weighted average:  $0.7 * 2 + 0.1 * 6 + 0.2 * 3 = 2.6$

**2.40.2** For a 25% improvement, we must reduce the CPU to  $2.6 * .75 = 1.95$ . Thus, we want  $0.7 * x + 0.1 * 6 + 0.2 * 3 < = 1.95$ . Solving for x shows that the arithmetic instructions must have a CPI of at most 1.07.

**2.40.3** For a 50% improvement, we must reduce the CPU to  $2.6 * .5 = 1.3$ . Thus, we want  $0.7 * x + 0.1 * 6 + 0.2 * 3 < = 1.3$ . Solving for x shows that the arithmetic instructions must have a CPI of at most 0.14

```
2.41 lwr    $t1, $s0($s6), 2    # t1 = A[f]
      add    $s0, $s0, 1
      lwr    $t2, $s0($s6), 2    # t2 = A[1+f]
      add    $t0, $t1, $t2      # t0 = A[f] + A[1+f]
      swr    $t0, $s1($s7), 2    # B[g] = A[f] + A[1+f] (don't
                                # need scaled store here)
```

```
2.42 lwr    $t0, $s3($s6),2     # $t0 <-- A[i]
      lwr    $t1, $s4($s6),2     # $t1 <-- A[j]
      add    $t0, $t0,$t1        # t1 <-- A[i] + A[j]
      sw     $t0,32($s7)         # B[8] <-- A[i] + A[j]
```

